

Deep Learning Predictive Modeling in Finance | CNNs

aiquants Research | © Dayton Nyamai

Application of CNNs in Market Direction Prediction

Summary

This project centers around the utilization of Convolutional Neural Network models in deep learning for financial market analysis. It entails the development of a Python code that leverages historical data to forecast the price movement direction of a financial instrument. The code encompasses several stages, including data preprocessing, feature engineering, and visualization, all tailored specifically for the field of python-finance

```
[124]: # Import the necessary libraries
import pandas as pd
import numpy as np
```

```
[125]: from pylab import mpl, plt
plt.style.use('seaborn-v0_8')
mpl.rcParams['font.family'] = 'serif'
%matplotlib inline
```

The Data

```
[126]: # Load the historical data and drop any row with missing values
url = 'https://raw.githubusercontent.com/dayton-nyamai/MarketDLModels/main/data/
↳data.csv'

raw = pd.read_csv(url, index_col=0, parse_dates=True).dropna()
#raw.info()
```

Select the symbol and create a DataFrame

```
[127]: symbol = ['EUR=']
data = pd.DataFrame(raw[symbol])
```

Align dates and rename the column containing the price data to 'price'.

```
[128]: start_date = data.index.min()
end_date = data.index.max()
data = data.loc[start_date:end_date]
data.rename(columns={'EUR': 'price'}, inplace=True)
data.round(4).head()
```

```
[128]:           price
Date
2010-01-04  1.4411
2010-01-05  1.4368
2010-01-06  1.4412
2010-01-07  1.4318
2010-01-08  1.4412
```

Calculate log returns and create direction column

```
[129]: data['returns'] = np.log(data['price'] / data['price'].shift(1))
data.dropna(inplace=True)
data['direction'] = np.where(data['returns'] > 0, 1, 0)
data.round(4).head()
```

```
[129]:           price  returns  direction
Date
2010-01-05  1.4368  -0.0030           0
2010-01-06  1.4412   0.0031           1
2010-01-07  1.4318  -0.0065           0
2010-01-08  1.4412   0.0065           1
2010-01-11  1.4513   0.0070           1
```

A histogram providing visual representation of the EUR log returns distribution

```
[130]: data['returns'].hist(bins=35, figsize=(10, 6));
# Add figure caption
plt.figtext(0.5, -0.01, 'Fig. 1.1 A histogram showing the distribution of EUR_
↳log returns ', style='italic',ha='center')

# Show the plot
plt.show()
```




Fig. 1.1 A histogram showing the distribution of EUR log returns

Second, create the features data by lagging the log returns and visualize it in combination with the returns data. We can use various visualization techniques such as scatter plots or line plots to compare the lagged log returns with the returns data.

Create lagged columns

```
[131]: lags = 5

cols = [ ]
for lag in range(1, lags+1):
    col = f'lag_{lag}'
    data[col] = data['returns'].shift(lag)
    cols.append(col)
data.dropna(inplace=True)

data.round(4).tail()
```

```
[131]:
```

	price	returns	direction	lag_1	lag_2	lag_3	lag_4	lag_5
Date								
2019-12-24	1.1087	0.0001	1	0.0007	-0.0038	0.0008	-0.0034	0.0006
2019-12-26	1.1096	0.0008	1	0.0001	0.0007	-0.0038	0.0008	-0.0034
2019-12-27	1.1175	0.0071	1	0.0008	0.0001	0.0007	-0.0038	0.0008
2019-12-30	1.1197	0.0020	1	0.0071	0.0008	0.0001	0.0007	-0.0038
2019-12-31	1.1210	0.0012	1	0.0020	0.0071	0.0008	0.0001	0.0007

Scatter plot based on features and labels data

```
[132]: data.plot.scatter(x='lag_1', y='lag_2', c='returns',
                        cmap='coolwarm', figsize=(10, 6), colorbar=True)

# Adding vertical and horizontal lines at 0
plt.axvline(0, c='r', ls='--')
plt.axhline(0, c='r', ls='--')

# Add figure caption
plt.figtext(0.4, -0.03,
            'Fig. 1.2 A scatter plot based on features and labels data',
            style='italic', ha='center')

# Show the plot
plt.show()
```




Fig. 1.2 A scatter plot based on features and labels data

With the dataset fully prepared, various deep learning techniques can be employed to forecast market movements based on the provided features. Additionally, these predictions can be utilized to rigorously backtest a trading strategy.

Deep Learning Models: Convolutional Neural Networks

Convolutional Neural Networks (CNNs) are a type of deep learning algorithm commonly used for image recognition tasks. However, they can also be applied to market movement predictions by treating the financial data as an image. CNNs excel at capturing spatial patterns and dependencies, making them suitable for analyzing the temporal patterns in market data. By training a CNN on historical market data and corresponding movement labels, it can learn to make predictions on unseen data, helping traders and investors make informed decisions.

Summary

In this task, we will create a CNNs model for predicting future market movements. We will also utilize the TPU (Tensor Processing Unit) VM cloud infrastructure from Google, for efficient training and inference.

Import the necessary libraries, tensorflow and its submodules

```
[133]: from sklearn.preprocessing import StandardScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten, Dense
from tensorflow.keras.optimizers import Adam
```

Split the data into training and test sets

```
[134]: cutoff = '2017-12-31'
training_data = data[data.index < cutoff].copy()
test_data = data[data.index >= cutoff].copy()
```

Standardize the training and test data.

```
[135]: mu, std = training_data.mean(), training_data.std()
training_data_ = (training_data - mu) / std
test_data_ = (test_data - mu) / std
```

Reshape the training and test data for CNN input

```
[136]: X_train = np.array(training_data_[cols])
X_train = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
y_train = np.array(training_data_['direction'])

X_test = np.array(test_data_[cols])
X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
y_test = np.array(test_data_['direction'])
```

Build the CNN model

```
[137]: # Build the CNN model
model = Sequential()
model.add(Conv1D(filters=32, kernel_size=3, activation='relu',
    ↪input_shape=(lags, 1)))
model.add(MaxPooling1D(pool_size=2))
```

```
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
```

Compile the model

```
[138]: optimizer = Adam(learning_rate=0.0001)
model.compile(optimizer=optimizer, loss='binary_crossentropy',
↳metrics=['accuracy'])
```

Train the model

```
[139]: model.fit(X_train, y_train, epochs=100, verbose=False, validation_split=0.2,
↳shuffle=False)

res = pd.DataFrame(model.history.history)
res[['accuracy', 'val_accuracy']].plot(figsize=(10, 6), style='--');
plt.figtext(0.5, 0.02, 'Fig. 1.3 Accuracy of the CNN model on training and
↳validation data per training step', style='italic',ha='center')
plt.show()
```




Fig. 1.3 Accuracy of the CNN model on training and validation data per training step

Evaluate the performance of the model on training data

```
[140]: train_loss, train_accuracy = model.evaluate(X_train, y_train)
```

```
63/63 [=====] - 0s 1ms/step - loss: 0.6856 - accuracy: 0.5556
```

Analyzing Accuracy in a Machine Learning Model: The accuracy metrics are essential in evaluating the performance of a machine learning model. Let's break down the metrics and understand their functionalities.

1. **Accuracy / train_accuracy:** The accuracy metric measures the proportion of correctly predicted instances out of the total number of instances. It provides an overall assessment of the model's performance. In our case, the accuracy value is 0.5556, indicating that the model correctly predicted 54.61% of the instances.
2. **Val_Accuracy:** The val_accuracy metric, also known as validation accuracy, measures the accuracy of the model on a separate validation dataset. It helps assess the model's generalization capability. In our case, the val_accuracy value is not explicitly mentioned in the output. However, it is expected to be plotted alongside the accuracy metric in the generated plot.
3. **train_loss:** The loss value calculated during the evaluation of the model on the training dataset. Loss is a measure of how well the model is performing in terms of the discrepancy between the predicted and actual values. In this case, the value of train_loss is 0.6856, suggesting that the model's predictions have a relatively high discrepancy from the actual values in the training dataset. A lower loss value indicates better model performance, so this relatively high loss value implies that the model may not be performing optimally on the training data.

Ideally, we want both accuracy and val_accuracy to be high. A high accuracy indicates that the model has learned the patterns in the training data well, while a high val_accuracy suggests that the model can generalize well to new data.

Insights on Accuracy Metrics

Let's discuss the insights we can gain from analyzing these metrics.

- **Trend:** The plot shows the trend of accuracy and val_accuracy over the epochs. We can analyze whether these metrics are improving, plateauing, or deteriorating over time. A consistent increase in both metrics indicates that the model is learning and improving its performance.
- **Overfitting:** If the accuracy metric keeps improving while the val_accuracy metric starts to plateau or decrease, it suggests that the model is overfitting. Overfitting occurs when the model becomes too specialized in the training data and fails to generalize well on unseen data.
- **Underfitting:** On the other hand, if both accuracy and val_accuracy remain low or do not show significant improvement, it indicates underfitting. Underfitting occurs when the model fails to capture the underlying patterns in the data and performs poorly on both training and validation datasets.
- **Convergence:** If both accuracy and val_accuracy reach a stable value and remain constant over the epochs, it suggests that the model has converged. Convergence indicates that the model has learned the patterns in the data and is not likely to improve further.

Analyzing the accuracy metrics help us understand the performance and behavior of a machine learning model. It allows us to make informed decisions regarding model optimization, such as

adjusting hyperparameters, increasing training data, or implementing regularization techniques.

Training Data

Make predictions on the training data

```
[141]: train_predictions = np.where(model.predict(X_train) > 0.5, 1, 0)
```

```
63/63 [=====] - 0s 1ms/step
```

Transforms the predictions into long-short positions, +1 and -1

```
[142]: training_data['prediction'] = np.where(train_predictions > 0, 1, -1)
```

The number of the resulting short and long positions, respectively.

```
[143]: training_data['prediction'].value_counts()
```

```
[143]: -1    1128
       1     879
       Name: prediction, dtype: int64
```

Trading Rules

In the benchmark case, i.e. `training_data['returns']`, we adopt a long position on the asset throughout the entire period. This means that we hold one unit of the asset for the entire duration. On the other hand, in the case of the CNNs strategy, i.e. `training_data['strategy']`, we take either a long or short position on the asset, i.e. one unit of the asset.

Calculates the strategy returns given the positions

```
[144]: training_data['strategy'] = training_data['prediction'] *
      ↪ training_data['returns']
      training_data[['returns', 'strategy']].sum().apply(np.exp)
```

```
[144]: returns    0.826569
      strategy    2.061973
      dtype: float64
```

Plots and compares the strategy performance to the benchmark performance (in-sample)

```
[145]: training_data[['returns', 'strategy']].cumsum().apply(np.exp).plot(figsize=(10,
      ↪ 6));
      plt.figtext(0.5, 0.05, 'Fig. 1.4 Gross performance of EUR/USD compared to the
      ↪ CNNs-based strategy', style='italic', ha='center')
      plt.figtext(0.5, -0.01, '(in-sample, no transaction costs)',
      ↪ style='italic', ha='center')

      plt.show()
```




Fig. 1.4 Gross performance of EUR/USD compared to the CNNs-based strategy
(in-sample, no transaction costs)

The strategy exhibits a modestly superior performance compared to the passive benchmark case on the training data set (in-sample, excluding transaction costs). However, the true measure of its effectiveness lies in its out-of-sample performance on the test data set, which is of utmost interest.

Testing Data

Evaluate the performance of the model on testing data

```
[146]: model.evaluate(X_test, y_test)
```

```
16/16 [=====] - 0s 1ms/step - loss: 0.6918 - accuracy: 0.5249
```

```
[146]: [0.6918313503265381, 0.5248509049415588]
```

Make predictions on the test data

```
[147]: test_predictions = np.where(model.predict(X_test) > 0.5, 1, 0)

# Transforms the predictions into long-short positions, +1 and -1
test_data['prediction'] = np.where(test_predictions > 0, 1, -1)
```

```
16/16 [=====] - 0s 1ms/step
```

The number of the resulting short and long positions, respectively.

```
[148]: test_data['prediction'].value_counts()
```

```
[148]: -1    346  
      1    157  
      Name: prediction, dtype: int64
```


Calculate the strategy returns given the positions

```
[149]: test_data['strategy'] = test_data['prediction'] * test_data['returns']  
      test_data[['returns', 'strategy']].sum().apply(np.exp)
```

```
[149]: returns    0.934478  
      strategy    1.122643  
      dtype: float64
```

Plots and compares the strategy performance to the benchmark performance (out-of-sample)

```
[150]: test_data[['returns', 'strategy']].cumsum().apply(np.exp).plot(figsize=(10, 6));  
      plt.figtext(0.5, 0.05, 'Fig. 1.5 Gross performance of EUR/USD compared to the_  
      ↪DNNs-based strategy', style='italic',ha='center')  
      plt.figtext(0.5, -0.01, '(out-of-sample, no transaction costs)',_  
      ↪style='italic',ha='center')  
  
      plt.show()
```



*Fig. 1.5 Gross performance of EUR/USD compared to the DNNs-based strategy
(out-of-sample, no transaction costs)*

Adding Different Types of Features

To this end, the analysis has primarily focused on the use of log returns. It is, of course, possible to not only incorporate more classes or categories, but also add other types of features to the mix, such as volatility, momentum, distance measures, etc – the sky is the limit. In our next notebook we will graft all these features into the data.